

RAG PDF Brief

- Retrieval-Augmented Generation combines search with LLM generation.
- Pipeline steps: ingest documents, chunk text, build embeddings, index vectors.
- At query time, retrieve top-k chunks and pass them to the model as context.
- Metadata fields are critical: source, title, section, chunk_id.

Chunking Tips

- Fixed-size chunks are predictable and simple.
- Structure-based chunks preserve semantic boundaries and headings.
- Compare both strategies using quality of retrieval and context coherence.